

Efficient Computation of Expectations under Spanning Tree Distributions

September 8, 2021

Dependency Trees

- ▶ Classical representation of text
- ▶ Similar to phrase structure grammars
 - ▶ Phrase structure: Labels groups of words
 - ▶ Dependency: Labels relationship between pairs of words
- ▶ Useful if a language has free word order
- ▶ Spanning trees




Figure: Example of a dependency tree

Efficient Computation of Expectations under Spanning Tree Distributions

- ▶ A framework for computing expectations (of decomposable functions) over spanning trees
 - ▶ Unifies previous algorithms
 - ▶ Additionally show large asymptotic and empirical speed improvements over particular past implementations
- ▶ Relies on connections between moments and derivatives
 - ▶ Uses automatic differentiation for easy to implement and efficient algorithms
- ▶ Fun demonstration of a general ‘inference with automatic differentiation’ recipe
 - ▶ Compute partition function
 - ▶ Use AD to compute expectations

Outline

Goal is to compute expectations over spanning trees

- ▶ Background: Distributions over (spanning) trees
- ▶ Method: Connecting expectations to derivatives
 - ▶ Use properties of our choice of tree distribution and decomposable functions
 - ▶ Stitch together into efficient algorithms with automatic differentiation
- ▶ Computational complexity

Distributions over Trees

Assuming fixed length sentence, with N nodes:

- ▶ Weighted edges

$$(i \xrightarrow{w_{ij}} j) \in \mathcal{E}$$

- ▶ Trees weights

$$w(d) := \prod_{(i \rightarrow j) \in d} w_{ij}$$

- ▶ Tree probability obtained via normalization

$$p(d) := \frac{w(d)}{Z},$$

where

$$Z := \sum_{d \in \mathcal{D}} w(d) = \sum_{d \in \mathcal{D}} \prod_{(i \rightarrow j) \in d} w_{ij}$$

- ▶ Computation of Z is tractable despite exponentially many trees in \mathcal{D}

Distributions over spanning trees




Figure: Example of a dependency tree

Matrix-Tree Theorem (MTT)

- ▶ Compute partition function Z using the graph Laplacian $L \in \mathbb{R}^{N \times N}$: $Z = |L|$.
 - ▶ Simple (undirected, unweights) graphs: $L = D - A$, where D is the degree matrix and A the adjacency
 - ▶ Weighted directed graphs:

$$L_{ij} := \begin{cases} \sum_{i' \in \mathcal{N} \setminus \{j\}} w_{i'j} & \text{if } i = j \\ -w_{ij} & \text{otherwise} \end{cases}$$

- ▶ Determinant can be computed in $O(N^3)$ time
- ▶ Flexibility in choice of L depending on example, care is needed
- ▶ Results are general for choice of Laplacian L

Expectations

- ▶ Main goal is to compute expectations / totals

$$\mathbb{E}_d[f(d)] := \sum_{d \in \mathcal{D}} p(d)f(d) = \frac{1}{Z} \sum_{d \in \mathcal{D}} w(d)f(d) = \frac{1}{Z} \bar{f}$$

- ▶ Need to consider every unique tree with nonzero mass unless $f : \mathcal{D} \rightarrow \mathbb{R}^F$ decomposes
- ▶ Consider two families of decomposable f
 - ▶ Additively decomposable: Shannon entropy, KL
 - ▶ Second-order additively decomposable: Gradient of entropy, covariance

Intuition for decompositions

- ▶ Recall: Tree distribution is a distribution over a binary vector of size $O(N^2)$
- ▶ Compute expectations by
 - ▶ Storing the relevant parts of trees and their marginal probabilities
 - ▶ Applying f to those parts
- ▶ Consider different levels of decomposability of f
 - ▶ Does not decompose: exponential num trees
 - ▶ Decomposes over pairs of edges: $O(N^4)$ evals
 - ▶ Decomposes over edges: $O(N^2)$ evals
- ▶ Combine $p(\text{part})f(\text{part})$ to compute expectation (unnormalized total using $p(\text{part}) = \frac{1}{Z} \tilde{w}_{\text{part}}$)

(Unnormalized) marginals

- ▶ Unary marginals (prob of one edge appearing in a tree)

$$p((i \rightarrow j) \in d) = \frac{1}{Z} \sum_{d \in \mathcal{D}_{ij}} w(d) = \frac{1}{Z} \widetilde{w}_{ij}$$

- ▶ Binary marginals (prob of two edges appearing together in a tree)

$$p((i \rightarrow j), (k \rightarrow l) \in d) = \frac{1}{Z} \sum_{d \in \mathcal{D}_{ij,kl}} w(d) = \frac{1}{Z} \widetilde{w}_{ij,kl}$$

- ▶ Want to show that
 - ▶ Can decompose totals into sums over marginals (enumerate parts of trees)
 - ▶ Can compute marginals using automatic differentiation (cheap gradient principle, simple implementation)
 - ▶ Avoid storing full probability tables (requires low-level tricks, not our focus / see paper for details)

Additively decomposable functions

Allows us to decompose $r : \mathcal{D} \rightarrow \mathbb{R}^R$ into functions of edges and combine with (unnormalized) unary marginals (over those tree edges)

$$r(d) = \sum_{(i \rightarrow j) \in d} r_{ij}, \quad (1)$$

with $r_{ij} \in \mathbb{R}^R$

Additively decomposable functions: Marginals

Marginals are expectations

$$p((i \rightarrow j) \in d) = \mathbb{E}_{p(d)} [1((i \rightarrow j) \in d)]_{ij} = [\sum_d p(d)r(d)]_{ij}$$

Set

$$r(d) = 1((i \rightarrow j) \in d)$$

Proof

$$\begin{aligned} 1((i \rightarrow j) \in d) &= \sum_{(k \rightarrow l) \in d} 1(kl = ij) \\ &\Rightarrow r_{ij} = 1(kl = ij) \end{aligned}$$

Additively decomposable functions: Shannon Entropy

Shannon entropy is an expectation

$$\mathbb{E}_{p(d)} [-\log p(d)]$$

Set

$$r(d) = -\log p(d)$$

Proof

$$\begin{aligned} -\log p(d) &= -\log\left(\frac{1}{Z} \prod_{(i \rightarrow j) \in d} w_{ij}\right) \\ &= \log Z - \sum_{(i \rightarrow j) \in d} \log w_{ij} \\ \Rightarrow r_{ij} &= \frac{1}{N} \log Z - \log w_{ij} \end{aligned}$$

First order totals: Recipe

- ▶ Rewrite unary marginal as gradient
- ▶ Rewrite total as sum over unary parts

First order totals: Marginal

For any edge $(i \rightarrow j)$,

$$\widetilde{w}_{ij} := \sum_{d \in \mathcal{D}_{ij}} w(d) = \frac{\partial Z}{\partial w_{ij}} w_{ij}, \quad (2)$$

where $\mathcal{D}_{ij} := \{d \in \mathcal{D} \mid (i \rightarrow j) \in d\}$

- Significantly easier to implement via AD than manually

First order totals: Marginal derivation

$$\widetilde{w}_{ij} = \sum_{d \in \mathcal{D}_{ij}} w(d) \quad (\text{Defn})$$

$$= \sum_{d \in \mathcal{D}_{ij}} \prod_{(i' \rightarrow j') \in d} w_{i'j'} \quad (\text{Defn of } w(d))$$

$$= w_{ij} \sum_{d \in \mathcal{D}_{ij}} \prod_{\substack{(i' \rightarrow j') \in \\ d \setminus \{(i \rightarrow j)\}}} w_{i'j'} \quad (\text{Pull out } w_{ij})$$

$$= w_{ij} \frac{\partial}{\partial w_{ij}} \sum_{d \in \mathcal{D}} \prod_{(i' \rightarrow j') \in d} w_{i'j'} \quad (\text{Sum-prod deriv trick})$$

$$= \frac{\partial Z}{\partial w_{ij}} w_{ij} \quad (\text{Defn of } Z)$$

First order totals: Total

For any additively decomposable function $r: \mathcal{D} \mapsto \mathbb{R}^R$,

$$\bar{r} = \sum_{(i \rightarrow j) \in \mathcal{E}} \widetilde{w}_{ij} r_{ij} = \sum_{(i \rightarrow j) \in \mathcal{E}} \frac{\partial Z}{\partial w_{ij}} w_{ij} r_{ij} \quad (3)$$

Proof.

$$\bar{r} = \sum_{d \in \mathcal{D}} w(d) r(d) \quad (\text{Defn})$$

$$= \sum_{d \in \mathcal{D}} w(d) \sum_{(i \rightarrow j) \in d} r_{ij} \quad (\text{Defn of } r(d))$$

$$= \sum_{d \in \mathcal{D}} \sum_{(i \rightarrow j) \in d} w(d) r_{ij} \quad (\text{Distributive prop})$$

$$= \sum_{(i \rightarrow j) \in \mathcal{E}} \sum_{d \in \mathcal{D}_{ij}} w(d) r_{ij} \quad (\text{Commutative prop})$$

$$= \sum_{(i \rightarrow j) \in \mathcal{E}} \widetilde{w}_{ij} r_{ij} \quad (\text{Defn of } \widetilde{w}_{ij})$$

First order totals: Algorithm

- 1: **def** $T_1(w: \mathcal{E} \mapsto \mathbb{R}, r: \mathcal{E} \mapsto \mathbb{R}^R)$:
- 2: \triangleright *Compute first-order total; requires $\mathcal{O}(N^3 R')$ time, $\mathcal{O}(N^2 + R)$ space.*
- 3: Compute all \widetilde{w}_{ij} via AD in $\mathcal{O}(N^3)^1$
- 4: $\bar{r} \leftarrow \sum_{(i \rightarrow j) \in \mathcal{E}} \widetilde{w}_{ij} r_{ij}$ $\triangleright \mathcal{O}(N^2 R')$
- 5: **return** \bar{r}

¹Constant factor of $|L|$ runtime with AD

Second-order additively decomposable functions

- ▶ Allows us to decompose into functions of pairs of edges and combine with (unnormalized) binary marginals

$$t(d) = r(d)s(d)^{\top} \quad (4)$$

- ▶ Outer product of two additively decomposable functions, $r: \mathcal{D} \mapsto \mathbb{R}^R$ and $s: \mathcal{D} \mapsto \mathbb{R}^S$
- ▶ $t(d) \in \mathbb{R}^{R \times S}$ is generally a matrix.

Second-order additively decomposable functions: Binary marginals

Binary marginals are expectations

$$\begin{aligned} p((i \rightarrow j), (k \rightarrow l) \in d) &= \mathbb{E}_{p(d)} [1((i \rightarrow j), (k \rightarrow l) \in d)]_{ij,kl} \\ &= \left[\sum_d p(d) r(d) s(d)^\top \right]_{ij,kl} \end{aligned}$$

Set

$$r(d) = 1((i \rightarrow j) \in d), \quad s(d) = 1((k \rightarrow l) \in d)$$

which we know are additively decomposable

Second-order totals: Recipe

- ▶ Rewrite binary marginal as Hessian of partition function
- ▶ Rewrite grad of intermediate total as Hessian-vector product
- ▶ Rewrite second-order total as Jacobian-matrix product or Hessian-matrix product

Second-order totals: Marginal

Unnormalized binary marginals

$$\widetilde{w_{ij,kl}} := \sum_{d \in \mathcal{D}_{ij,kl}} w(d) = \frac{\partial^2 Z}{\partial w_{ij} \partial w_{kl}} w_{ij} w_{kl} \quad (5)$$

where $\mathcal{D}_{ij,kl} := \{d \in \mathcal{D} \mid (i \rightarrow j) \in d, (k \rightarrow l) \in d\}$

Second-order totals: Marginal

Same tricks as first-order marginals

$$\begin{aligned}\widetilde{w_{ij,kl}} &= \sum_{d \in \mathcal{D}_{ij,kl}} w(d) \\ &= \sum_{d \in \mathcal{D}_{ij,kl}} \prod_{(k' \rightarrow l') \in d} w_{k'l'} \\ &= w_{ij} w_{kl} \frac{\partial^2}{\partial w_{ij} \partial w_{kl}} \sum_{d \in \mathcal{D}} \prod_{(i' \rightarrow j') \in d} w_{i'j'} \\ &= \frac{\partial^2 Z}{\partial w_{ij} \partial w_{kl}} w_{ij} w_{kl}\end{aligned}$$

Second-order totals: Grad of intermediate total

Write $\nabla \bar{r}$ in terms of Hessian $\nabla^2 Z$, as a step towards writing \bar{t} in terms of $\nabla^2 Z$

$$\begin{aligned} & w_{ij} \frac{\partial \bar{r}}{\partial w_{ij}} \\ &= w_{ij} \frac{\partial}{\partial w_{ij}} \left(\sum_{(k \rightarrow l) \in \mathcal{E}} \frac{\partial Z}{\partial w_{kl}} w_{kl} r_{kl} \right) \\ &= w_{ij} \frac{\partial Z}{\partial w_{ij}} r_{ij} + w_{ij} \sum_{(k \rightarrow l) \in \mathcal{E}} \frac{\partial^2 Z}{\partial w_{ij} \partial w_{kl}} w_{kl} r_{kl} \\ &= \widetilde{w}_{ij} r_{ij} + \sum_{(k \rightarrow l) \in \mathcal{E}} \widetilde{w}_{ij,kl} r_{kl} \end{aligned}$$

Second-order totals: Total

We can both

- ▶ Write total \bar{t} in terms of Jacobian $\frac{\partial \bar{r}}{\partial w_{ij}}$ or gradients $\nabla \bar{r}_n$
 - ▶ Run backward for each dimension or \bar{r} : $O(RN^3 + R^2N^2)$
- ▶ Write total \bar{t} in terms of Hessian $\nabla^2 Z$
 - ▶ Hessian is large $O(N^4)$ entries
 - ▶ This takes $O(N^4)$ time, but can be optimized to $O(N^2 + RS)$ space

Second-order totals: Total

$$\begin{aligned}\bar{t} &= \sum_{d \in \mathcal{D}} w(d)r(d)s(d)^\top \\ &= \sum_{d \in \mathcal{D}} w(d)r(d) \sum_{(i \rightarrow j) \in d} s_{ij}^\top \\ &= \sum_{d \in \mathcal{D}} \sum_{(i \rightarrow j) \in d} w(d)r(d)s_{ij}^\top \\ &= \sum_{(i \rightarrow j) \in \mathcal{E}} \sum_{d \in \mathcal{D}_{ij}} w(d)r(d)s_{ij}^\top \\ &= \sum_{(i \rightarrow j) \in \mathcal{E}} w_{ij} \frac{\partial}{\partial w_{ij}} \left(\sum_{d \in \mathcal{D}} w(d)r(d) \right) s_{ij}^\top \\ &= \sum_{(i \rightarrow j) \in \mathcal{E}} w_{ij} \frac{\partial \bar{r}}{\partial w_{ij}} s_{ij}^\top\end{aligned}$$

Second-order totals: Jacobian version

- 1: **def** $T_2^v(w: \mathcal{E} \mapsto \mathbb{R}, r: \mathcal{E} \mapsto \mathbb{R}^R, s: \mathcal{E} \mapsto \mathbb{R}^S) :$
- 2: \triangleright *Compute second-order total with gradient-vector products; requires $\mathcal{O}(R(N^3 + N^2 R' + N^2 S'))$ time, $\mathcal{O}(N^2 R + R S)$ space.*
- 3: **for** $n = 1 \dots R :$ $\triangleright \mathcal{O}(R(N^3 + N^2 R'))$
- 4: Compute $\nabla \bar{r}_n$ using reverse-mode AD on $[T_1(w, r)]_n$
- 5: \triangleright *Requires $\mathcal{O}(N^2 R S')$*
- 6: **return** $\sum_{(i \rightarrow j) \in \mathcal{E}} \frac{\partial \bar{r}}{\partial w_{ij}} w_{ij} s_{ij}^\top$

► Recommended approach

Second-order totals: Hessian version

- 1: **def** $T_2^h(w: \mathcal{E} \mapsto \mathbb{R}, r: \mathcal{E} \mapsto \mathbb{R}^R, s: \mathcal{E} \mapsto \mathbb{R}^S)$:
 - 2: ▷ *Compute second-order total by materializing Hessian; requires $\mathcal{O}(N^4 R' S')$ time, $\mathcal{O}(N^2 + R S)$ space.*
 - 3: Compute all \widetilde{w}_{ij} via ∇Z
 - 4: Compute all $\widetilde{w}_{ij,kl}$ via $\nabla^2 Z$ (can compute this efficiently)
 - 5: ▷ *Requires $\mathcal{O}(N^4 R' S')$*
 - 6: **return** $\sum_{(i \rightarrow j) \in \mathcal{E}} \widetilde{w}_{ij} r_{ij} s_{ij}^\top + \sum_{(k \rightarrow l) \in \mathcal{E}} \widetilde{w}_{ij,kl} r_{ij} s_{kl}^\top$
- ▶ Better than previous algo at dealing with sparsity, i.e. if R and S are really big and sparse
 - ▶ Can come up with a 3rd algo that has the best of both worlds by computing the second term involving the Hessian-matrix product more efficiently

- [1] Dougal Maclaurin, David Duvenaud, and Ryan Adams. Gradient-based hyperparameter optimization through reversible learning. In *International Conference on Machine Learning*, pages 2113–2122, 2015.
- [2] Amirreza Shaban, Ching-An Cheng, Nathan Hatch, and Byron Boots. Truncated back-propagation for bilevel optimization. In *International Conference on Artificial Intelligence and Statistics*, pages 1723–1732, 2019.