

Efficient Computation of Expectations under Spanning Tree Distributions

September 8, 2021

Dependency Trees

- ▶ Classical representation of text
- ▶ Similar to phrase structure grammars
 - ▶ Phrase structure: Labels groups of words
 - ▶ Dependency: Labels relationship between pairs of words
- ▶ Useful if a language has free word order
- ▶ Spanning trees




Figure: Example of a dependency tree

Efficient Computation of Expectations under Spanning Tree Distributions

- ▶ A framework for computing expectations (of decomposable functions) over spanning trees
 - ▶ Unifies previous algorithms
 - ▶ Additionally show large asymptotic and empirical speed improvements over particular past implementations
- ▶ Relies on connections between moments and derivatives
 - ▶ Uses automatic differentiation for easy to implement and efficient algorithms
- ▶ Fun demonstration of inference with automatic differentiation

Outline

Goal is to compute expectations over spanning trees

- ▶ Background: Distributions over (spanning) trees
- ▶ Method: Connecting expectations to derivatives
 - ▶ Use properties of **our** choice of tree distribution and decomposable functions
 - ▶ Stitch together into efficient algorithms with automatic differentiation
- ▶ Computational complexity

Distributions over Trees

Assuming fixed length sentence, with N nodes:

- ▶ Weighted edges

$$(i \xrightarrow{w_{ij}} j) \in \mathcal{E}$$

- ▶ Trees weights

$$w(d) := \prod_{(i \rightarrow j) \in d} w_{ij}$$

- ▶ Tree probability obtained via normalization

$$p(d) := \frac{w(d)}{Z},$$

where

$$Z := \sum_{d \in \mathcal{D}} w(d) = \sum_{d \in \mathcal{D}} \prod_{(i \rightarrow j) \in d} w_{ij}$$

- ▶ Computation of Z is tractable despite exponentially many trees in \mathcal{D}

Distributions over spanning trees




Figure: Example of a dependency tree

Matrix-Tree Theorem (MTT)

- ▶ Compute partition function Z using the graph Laplacian $L \in \mathbb{R}^{N \times N}$: $Z = |L|$.
 - ▶ Simple (undirected, unweights) graphs: $L = D - A$, where D is the degree matrix and A the adjacency
 - ▶ Weighted directed graphs:

$$L_{ij} := \begin{cases} \sum_{i' \in \mathcal{N} \setminus \{j\}} w_{i'j} & \text{if } i = j \\ -w_{ij} & \text{otherwise} \end{cases}$$

- ▶ Determinant can be computed in $O(N^3)$ time
- ▶ Flexibility in choice of L depending on example, care is needed
- ▶ Results are general for choice of Laplacian L

Expectations

- ▶ Main goal is to compute expectations / totals

$$\mathbb{E}_d[f(d)] := \sum_{d \in \mathcal{D}} p(d)f(d) = \frac{1}{Z} \sum_{d \in \mathcal{D}} w(d)f(d) = \frac{1}{Z} \bar{f}$$

- ▶ Need to consider every unique tree with nonzero mass unless $f : \mathcal{D} \rightarrow \mathbb{R}^F$ decomposes
- ▶ Consider two families of decomposable f
 - ▶ Additively decomposable: Shannon entropy, KL
 - ▶ Second-order additively decomposable: Gradient of entropy, covariance

Intuition for decompositions

- ▶ Recall: Tree distribution is a distribution over a binary vector of size $O(N^2)$
- ▶ Compute expectations by
 - ▶ Storing the relevant parts of trees and their marginal probabilities
 - ▶ Applying f to those parts
- ▶ Consider different levels of decomposability of f
 - ▶ Does not decompose: exponential num trees
 - ▶ Decomposes over pairs of edges: $O(N^4)$ evals
 - ▶ Decomposes over edges: $O(N^2)$ evals
- ▶ Combine $p(\text{part})f(\text{part})$ to compute expectation (unnormalized total using $p(\text{part}) = \frac{1}{Z} \tilde{w}_{\text{part}}$)

(Unnormalized) marginals

- Unary marginals (prob of one edge appearing in a tree)

$$p((i \rightarrow j) \in d) = \frac{1}{Z} \sum_{d \in \mathcal{D}_{ij}} w(d) = \frac{1}{Z} \widetilde{w}_{ij}$$

- Binary marginals (prob of two edges appearing together in a tree)

$$p((i \rightarrow j), (k \rightarrow l) \in d) = \frac{1}{Z} \sum_{d \in \mathcal{D}_{ij,kl}} w(d) = \frac{1}{Z} \widetilde{w}_{ij,kl}$$

- Want to show that
 - Can decompose totals into sums over marginals (enumerate parts of trees)
 - Can compute marginals using automatic differentiation (cheap gradient principle)
 - Avoid storing full probability tables (requires low-level tricks, not the focus / see paper for details)

Additively decomposable functions

- ▶ Allows us to decompose into functions of edges and combine with (unnormalized) unary marginals

First order totals

- ▶ Rewrite unary marginal as gradient
- ▶ Rewrwite total as sum over unary parts

First order totals: Marginal

▶ A

▶ B

First order totals: Total

▶ A

▶ B

First order totals: Algorithm

- ▶ Rewrite unary marginal as gradient of partition function
- ▶ Rewrwite first order total as sum over unary parts

Second-order additively decomposable functions

- ▶ Allows us to decompose into functions of pairs of edges and combine with (unnormalized) binary marginals

Second-order totals

- ▶ Rewrite binary marginal as Hessian of partition function
- ▶ Rewrite grad of intermediate total as Hessian-vector product
- ▶ Rewrite second-order total as Jacobian-matrix product or Hessian-matrix product

Second-order totals: Marginal

► blah

Second-order totals: Grad of intermediate total

► asdf

Second-order totals: Total

► asdf

- [1] Dougal Maclaurin, David Duvenaud, and Ryan Adams. Gradient-based hyperparameter optimization through reversible learning. In *International Conference on Machine Learning*, pages 2113–2122, 2015.
- [2] Amirreza Shaban, Ching-An Cheng, Nathan Hatch, and Byron Boots. Truncated back-propagation for bilevel optimization. In *International Conference on Artificial Intelligence and Statistics*, pages 1723–1732, 2019.